



KubeCon



CloudNativeCon

North America 2018

Using gRPC for Long-lived and Streaming RPCs

Eric Anderson, Staff Software Engineer, *Google*



Intended Audience



KubeCon



CloudNativeCon

North America 2018

Have used gRPC

- Made a service
- Wrote some clients

Interested in more advanced use-cases

Overview



KubeCon



CloudNativeCon

North America 2018

Long-lived RPCs

Streaming RPCs

- Half Duplex
- Full Duplex

Long-lived Streaming RPCs



Long-lived RPCs

Streaming RPCs

- Half Duplex
- Full Duplex

Long-lived Streaming RPCs

Long-lived RPCs



KubeCon



CloudNativeCon

North America 2018

RPCs that last minutes, hours, days

Long-lived RPC Use-cases

Reduce polling, reduced latency; "Hanging GET"

- Watches/notifications

Long-lived RPC Issues

Load balancing

- Uneven backend memory usage
- MAX_CONNECTION_AGE can accumulate connections

TCP disconnects will fail calls; may take time to be detected

Deadline not very useful*

Not great for operations that take a long time, as you need to "reconnect"

Long-lived RPC Improvements

Load balancing: Have server occasionally close RPC

- If using MAX_CONNECTION_AGE, can use [MAX_CONNECTION_AGE_GRACE](#) to auto-kill as a back-up

Time to detect disconnections: Client-side Keepalive

Since not generally latency-sensitive, enable [Wait-for-ready](#)

Overview



KubeCon



CloudNativeCon

North America 2018

Long-lived RPCs

Streaming RPCs

- Half Duplex
- Full Duplex

Long-lived Streaming RPCs

Streaming RPCs

Zero-to-many messages (instead of one)

Messages are ordered

Streaming is independent in each direction

Streaming RPCs (Unary)



KubeCon

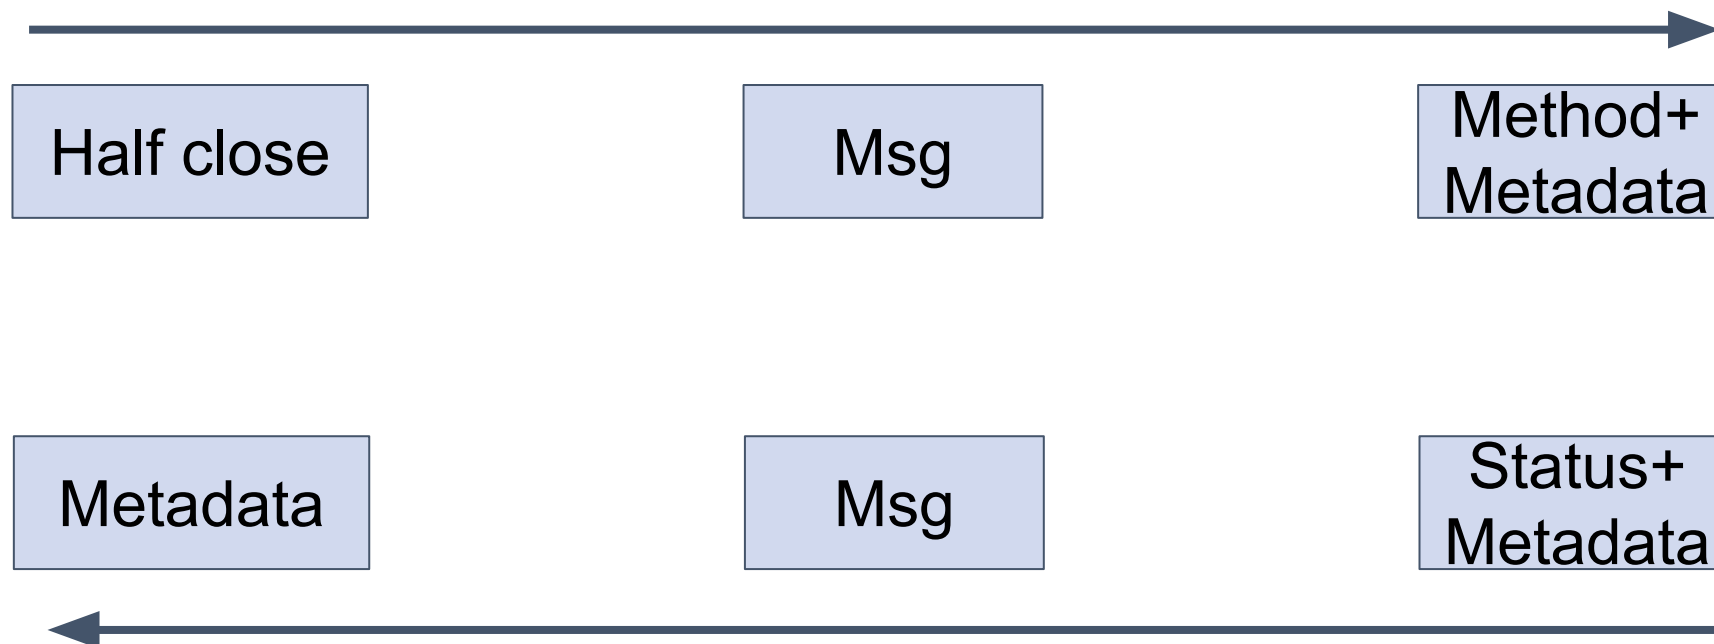


CloudNativeCon

North America 2018

Client

Server



Streaming RPCs (Streaming)



KubeCon

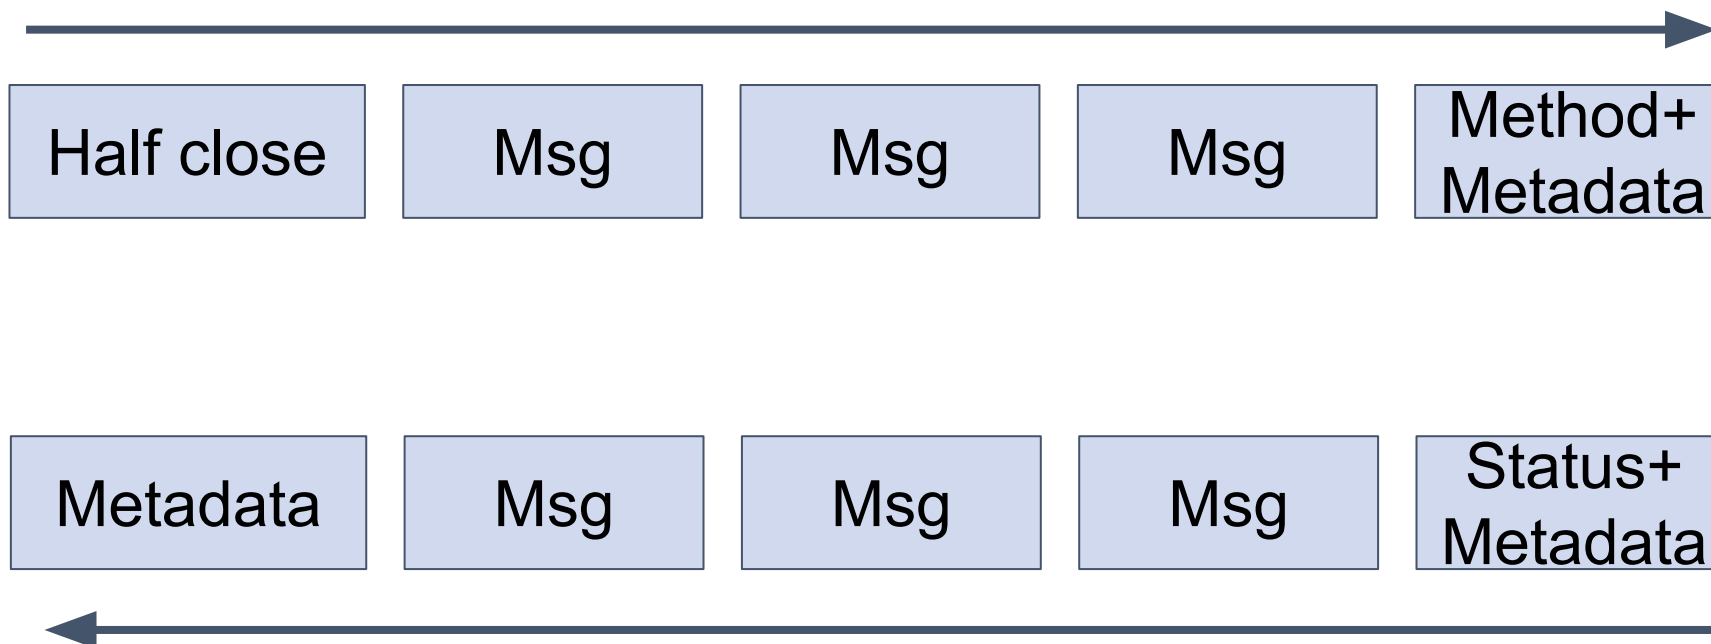


CloudNativeCon

North America 2018

Client

Server



Streaming RPCs

```
rpc UnaryCall (Request) returns (Response);
```

```
rpc ClientStreamingCall(stream Request) returns (Response);
```

```
rpc ServerStreamingCall (Request) returns (stream Response);
```

```
rpc BidirectionalCall(stream Request) returns (stream Response);
```

Streaming RPCs



KubeCon



CloudNativeCon

North America 2018

Bidirectional (Bidi) Streaming

- Half duplex. Client-streaming + Server-streaming
- Full duplex. More than one side can send at a time
 - Like TCP, but with messages instead of bytes (close semantics are a bit different, though)
 - No implicit acks; writes are only acked by responses



Long-lived RPCs

Streaming RPCs

- **Half Duplex**
- Full Duplex

Long-lived Streaming RPCs

Half-duplex Streaming Use-cases

Latency or memory reduction (e.g., speech to text)

- multiple small messages instead of a big message

Separation of response and “end of call” (e.g., watches)

Flow Control (“push-back”)

- Bulk uploads without needing to optimize chunk sizes
- Less “jerky” than one-at-a-time chunking (gives “pipelining”)

Half-duplex Streaming Use-cases



KubeCon



CloudNativeCon

North America 2018

Messages with state association

- Pinning to a backend
- Expands call lifetime (e.g., transactions)
- Reduced per-message setup cost (e.g., watches)
- Full-state followed by deltas (watches again...)

Half-duplex Streaming Issues

gRPC flow control may have large buffers (64 KB-4 MB)

gRPC flow control is point-to-point

Increased API complexity*

Server-streaming may require application-level retries*
(vs framework-level)

Tracing/stats muddled or missing

Half-duplex Streaming Improvements



KubeCon



CloudNativeCon

North America 2018

Flow control problems:

- use full duplex + application-level flow control

Tracing/stats: treating like unary could work okay



Long-lived RPCs

Streaming RPCs

- Half Duplex
- **Full Duplex**

Long-lived Streaming RPCs

Full-duplex Streaming Use-cases

TCP with messages

- Custom protocols

Application-level flow control (e.g., “messages,” “work items”)

Transactions

“Live” Reconfiguration

Bulk uploads, with reduced frequency of resumption

Use half-close to “hang up” instead of cancel

Full-duplex Streaming Issues



KubeCon



CloudNativeCon

North America 2018

Tracing/stats systems may be overly simplistic*

API/protocol complexity*

Involved application-level retry*

Flow-control-induced deadlocking

Lack of REST conversion*

Full-duplex Streaming Improvements



KubeCon



CloudNativeCon

North America 2018

Have at least one side be reading at any time

- If mixing control and data messages, use application-level flow control to limit memory usage

Overview



KubeCon



CloudNativeCon

North America 2018

Long-lived RPCs

Streaming RPCs

- Half Duplex
- Full Duplex

Long-lived Streaming RPCs

Long-lived Streaming Issues



KubeCon



CloudNativeCon

North America 2018

Load balancing (memory+cpu)

Tracing/stats systems may be overly simplistic*

Long-lived Streaming Improvements



KubeCon



CloudNativeCon

North America 2018

Load balancing: same as long-lived RPCs



GitHub/Gitter: @ejona86

Email: ejona@google.com (please CC mailing list)

Mailing List: grpc-io@googlegroups.com

Stack Overflow: #grpc #grpc-java