

# Extending Kubernetes LoadBalancer Using CRDs

**IBM Developer**

Wei Huang (IBM, [wei.huang1@ibm.com](mailto:wei.huang1@ibm.com))

Srini Brahmaroutu (IBM, [srbrahma@us.ibm.com](mailto:srbrahma@us.ibm.com))



- sig-scheduling



- sig-testing
- sig-storage

# Agenda

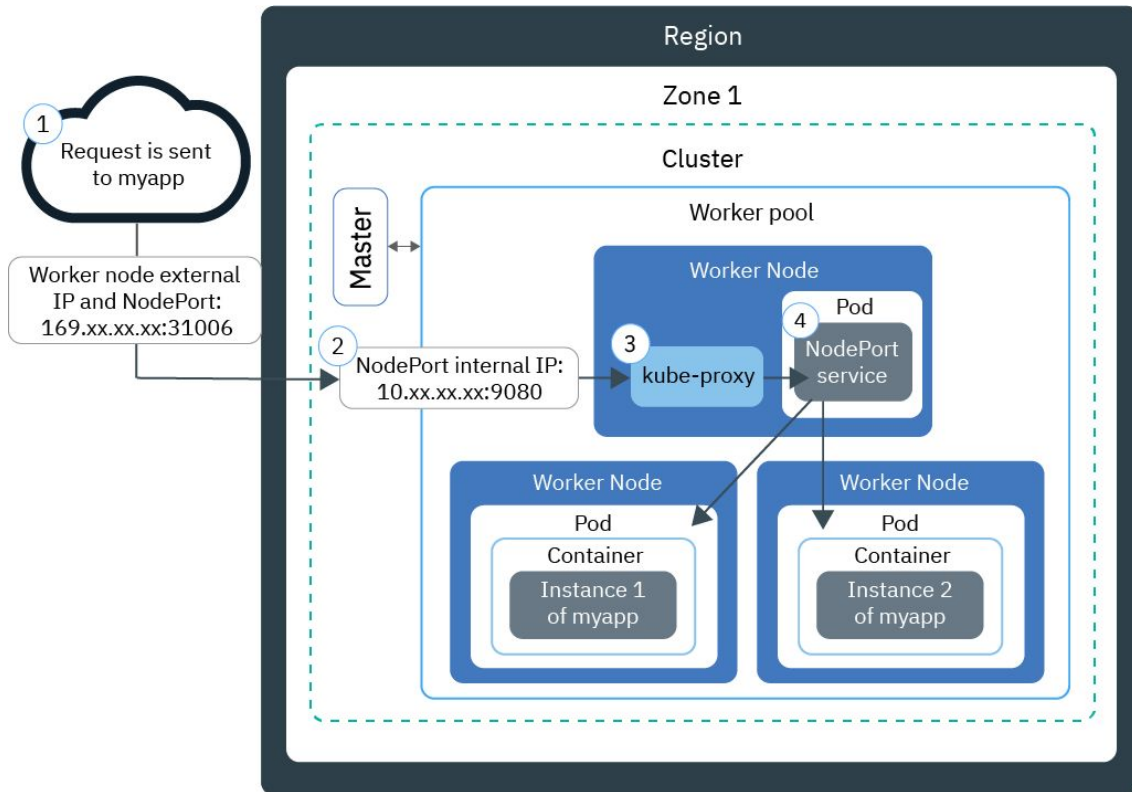
- How to Expose Kubernetes Workloads Externally
- Background / Motivations
- Shared LoadBalancer
- Demos
- Design / Implementation Details

# How to Expose Kubernetes Workloads Externally

# Kubernetes Basics - NodePort Service

## NodePort

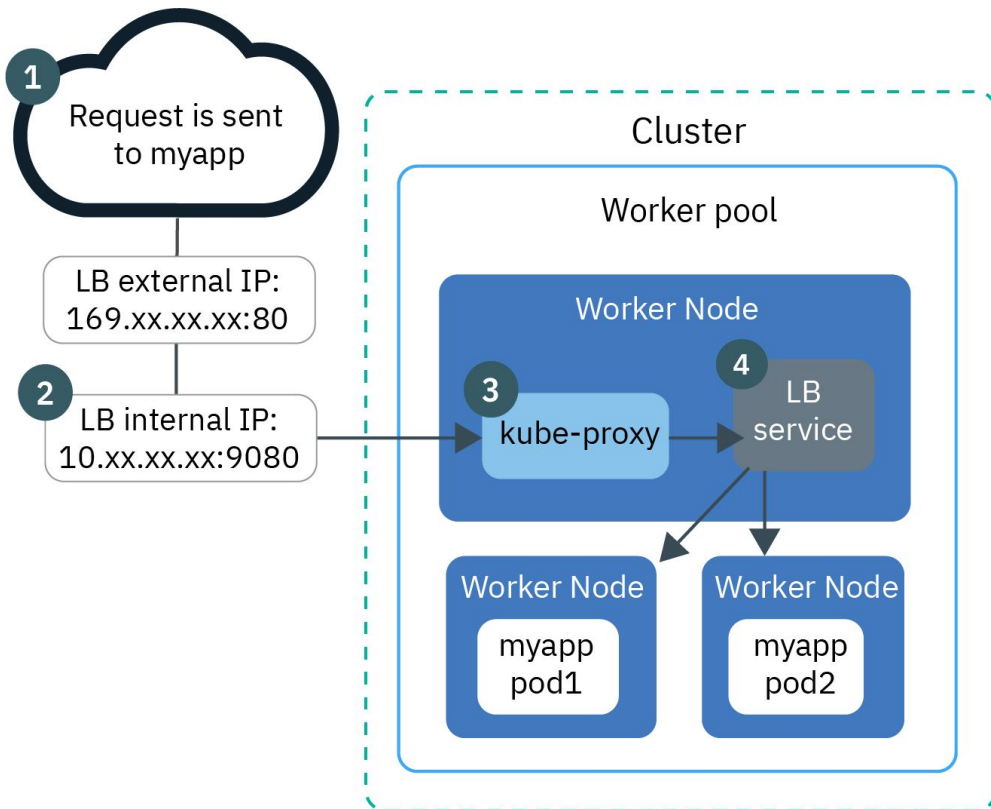
- Worker nodes need to have a public/external IP
- Ports opened on all worker nodes
- Ports range from 30000 to 32767



# Kubernetes Basics - LoadBalancer Service

## LoadBalancer

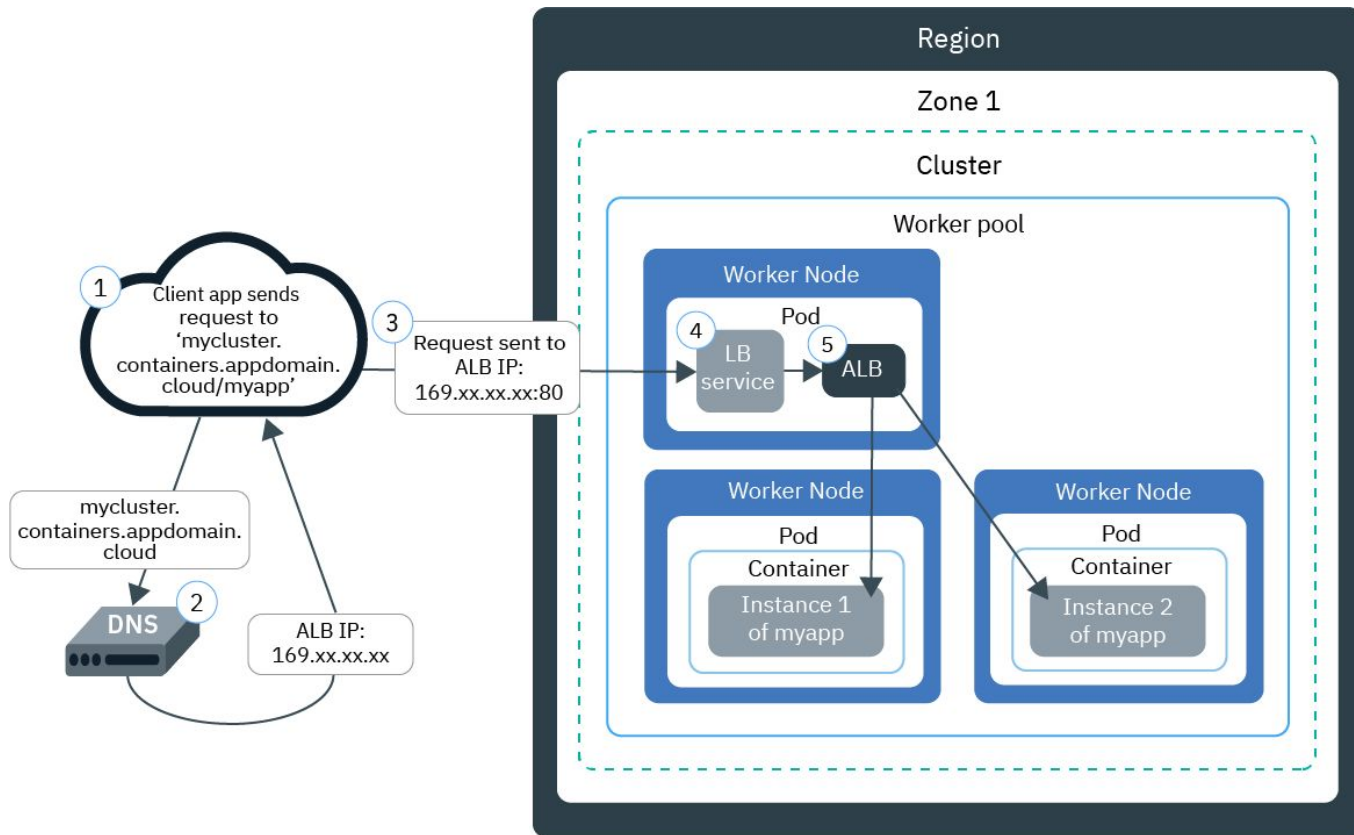
- EKS (Amazon)
- IKS (IBM)
- GKE (Google)
- AKS (Azure)
- ...



# Kubernetes Basics - Ingress

## Ingress

- L7
- Nginx/Envoy/...



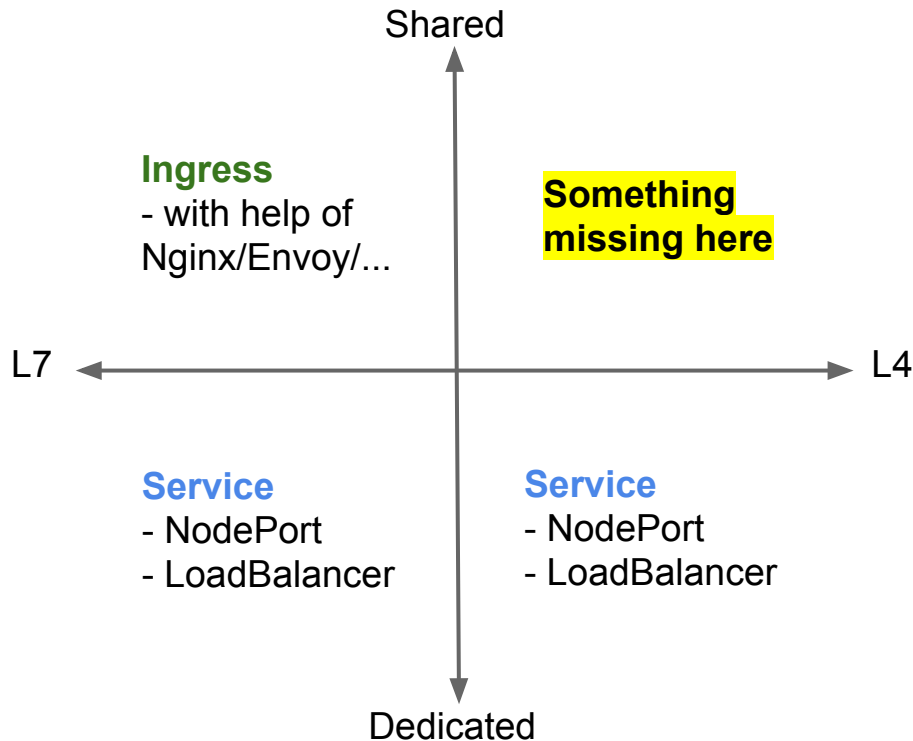
# Ways to Expose K8s Apps Externally

**Service** - for both L4/L7 traffic

- Type NodePort
- Type LoadBalancer

**Ingress** - for L7 traffic

- Shared via ingress controller



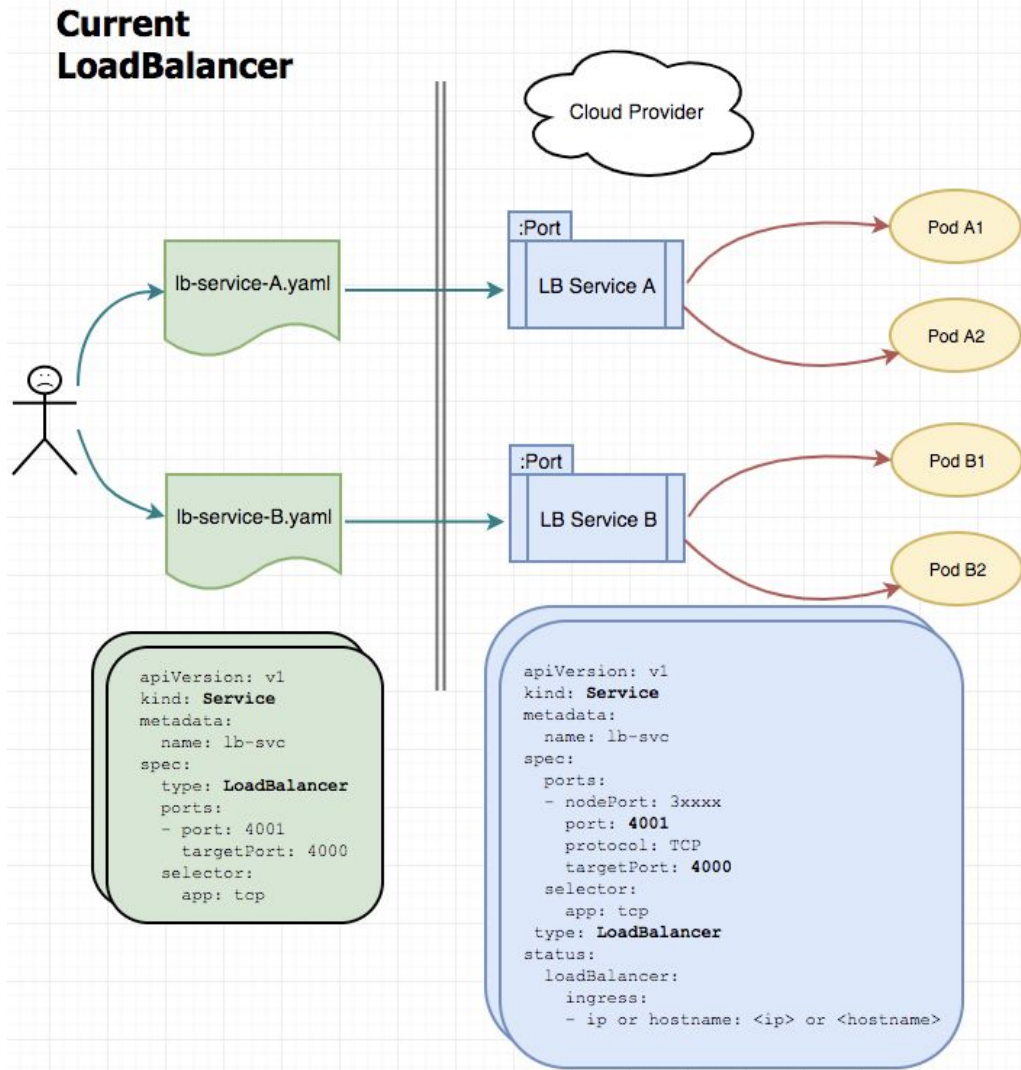


# Background / Motivations

# Background

An internal business requirement from an internal team.

- Two JDBC services (TCP)
- One data transferring service (UDP)
- One web console service (HTTP)



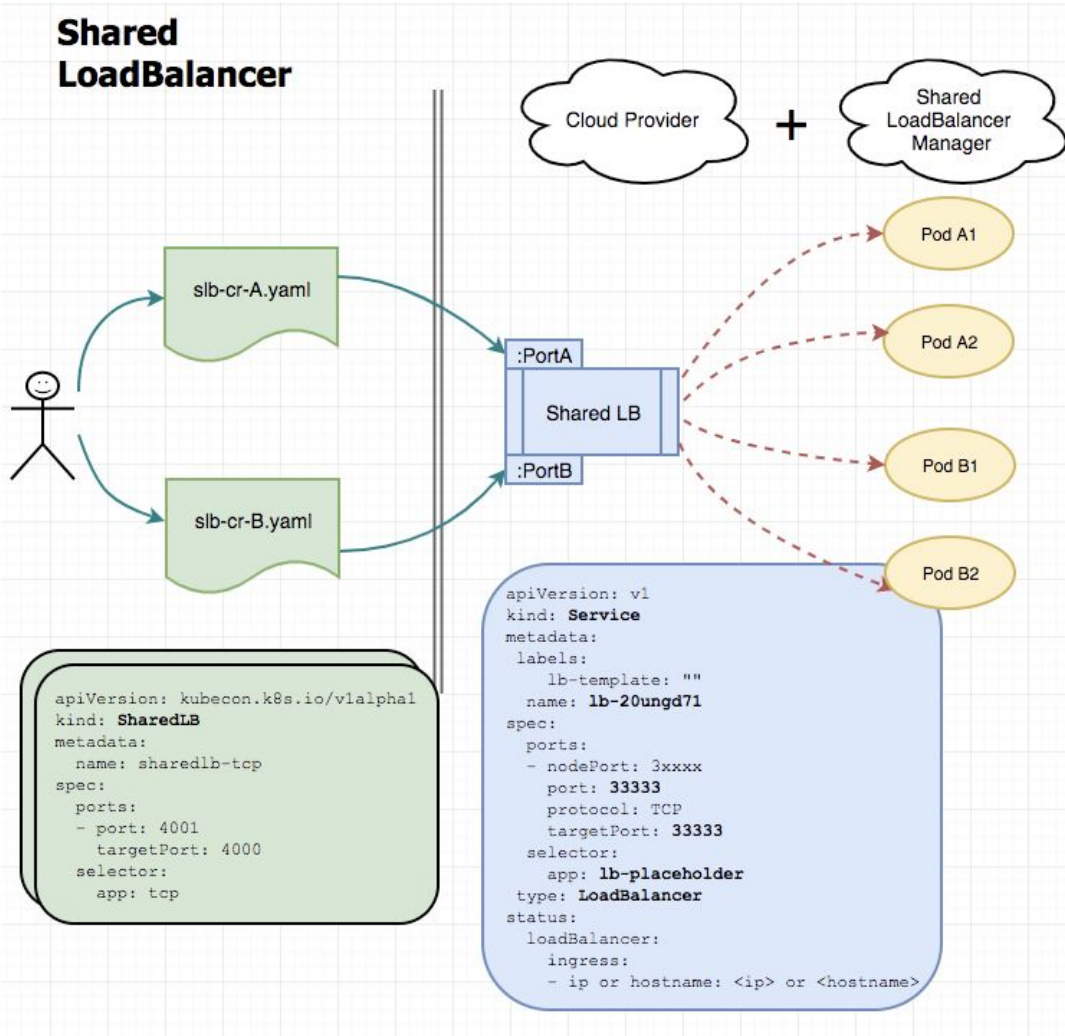
# Expected Goal

```
⇒ k create -f crs
sharedlb.kubecon.k8s.io/sharedlb-tcp1 created
sharedlb.kubecon.k8s.io/sharedlb-tcp2 created
sharedlb.kubecon.k8s.io/sharedlb-tcp3 created
sharedlb.kubecon.k8s.io/sharedlb-tcp4 created
wei.huang1@wei-mbp:~/gospace/src/github.com/Huang-Wei/shared-loadbalancer
⇒ k get slb
```

NAME	EXTERNAL-IP	PORT	PROTOCOL	REF
sharedlb-tcp1	169.62.88.170	4001	TCP	default/lb-z5lrv7he
sharedlb-tcp2	169.62.88.170	4002	TCP	default/lb-z5lrv7he
sharedlb-tcp3	169.62.88.170	4003	TCP	default/lb-z5lrv7he
sharedlb-tcp4	169.62.88.170	4004	TCP	default/lb-z5lrv7he

# Expected Goal (cont.)

```
➔ k create -f crs
sharedlb.kubecon.k8s.io/sharedlb-tcp1 created
sharedlb.kubecon.k8s.io/sharedlb-tcp2 created
sharedlb.kubecon.k8s.io/sharedlb-tcp3 created
sharedlb.kubecon.k8s.io/sharedlb-tcp4 created
wei.huang1@wei-mbp:~/gospace/src/github.com/Huang-Wei/shared-loadbalancer
➔ k get slb
NAME          EXTERNAL-IP    PORT    PROTOCOL  REF
sharedlb-tcp1 169.62.88.170  4001    TCP        default/lb-z5lrv7he
sharedlb-tcp2 169.62.88.170  4002    TCP        default/lb-z5lrv7he
sharedlb-tcp3 169.62.88.170  4003    TCP        default/lb-z5lrv7he
sharedlb-tcp4 169.62.88.170  4004    TCP        default/lb-z5lrv7he
```



# Motivations

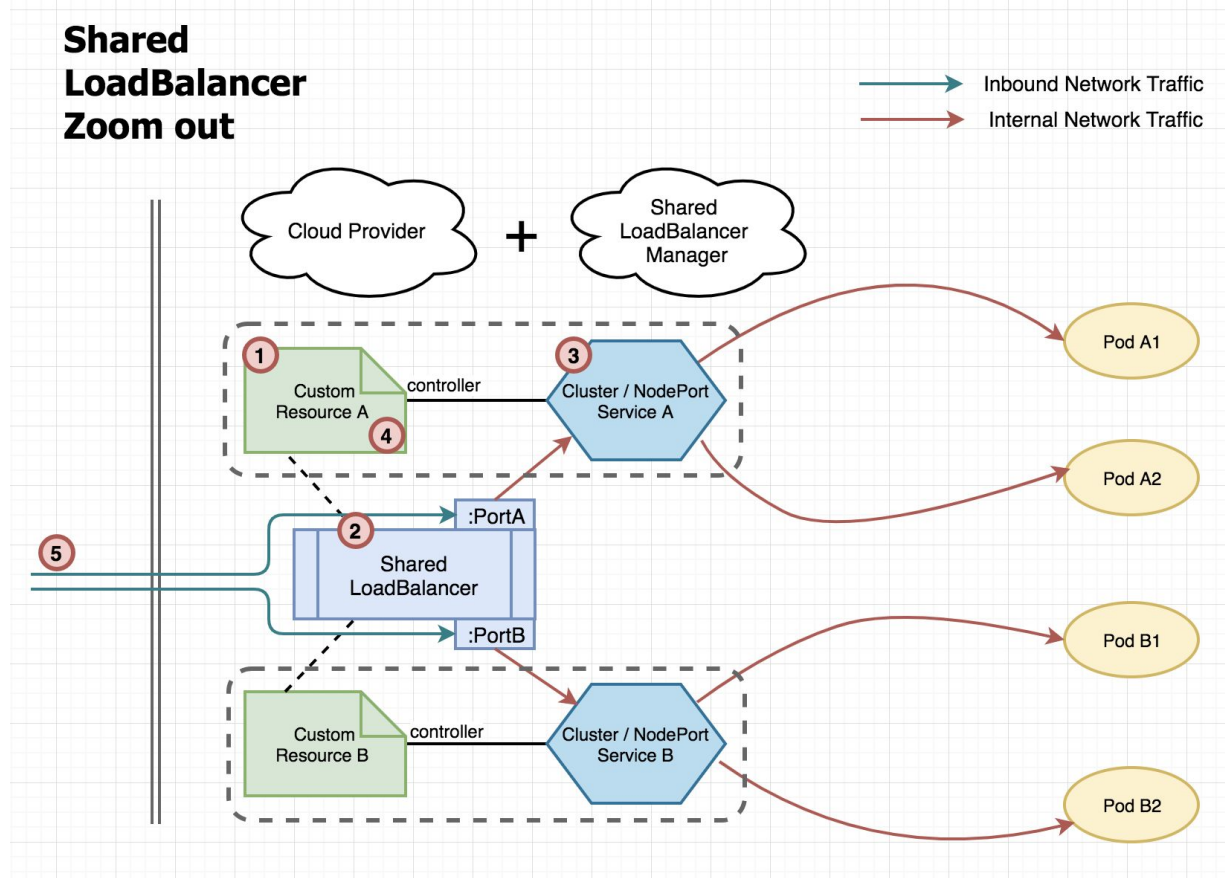
- Cost effective
- User friendly
- Reusing existing Kubernetes assets (don't reinvent wheel)
- Minimum operation efforts
- Consistent with Kubernetes roadmap

# Shared LoadBalancer (SLB)

# Problem Analysis

1. How to open **additional** ports (and firewall rules) on the “Shared” LoadBalancer
2. How to **associate** the ports with backing pods
3. How to give **accessing info** back to end-user

# Shared LoadBalancer Internals





# Demos

# Design / Implementation Details

# Design Considerations

1. Using CRD as the facade to end-user, ~~instead of Service with annotation.~~
2. Namespaced CRD ~~vs. Clustered CRD.~~
3. Create real LB on demand, ~~or prepare placeholder LBs in a pool.~~
4. Configure how many SharedLB CRs are mapped with one real LoadBalancer.
  - a. It can be a static way like 100 SharedLB CRs => 1 real LoadBalancer.
  - b. Or, can be done in a smart/dynamic way to take workload of each CR into consideration.
5. Support IKS, EKS, GKE and AKS

	EKS (Amazon)	IKS (IBM)	GKE (Google)	AKS (Azure)
<b>Core Extension Solution</b>	NodePort Service	Cluster Service with “externalIP”	NodePort Service	NodePort Service
<b>SDK Authentication</b>	aws_access_key_id aws_secret_access_key	APIKEY <i>(only needed when adding portable ip quota)</i>	oauth2 <i>(gcloud auth application-default login)</i>	Service principle and Role <i>(az ad sp create-for-rbac)</i>
<b>Forward Rule Firewall Rule</b>	Use SDK to operate	Auto Managed	Use SDK to operate	Use SDK to operate
<b>Accessing method</b>	<Hostname>:<Port>	<IP>:<Port>	<IP>:<Port>	<IP>:<Port>
<b>Limitations</b>	UDP not supported Latest version is 1.10	N/A	Incoming port limited to 30000~32767	N/A

# Future Considerations

1. Get feedback (sig-cloudprovider, sig-network, internal & external users)
2. Support on more cloudproviders
3. Testings (unit tests, integration tests, e2e tests) and CI
4. Features (essentially it's a scheduling problem)
  - a. Ports (completed)
  - b. Request/Limits
  - c. LeastRequested vs. MostRequested
  - d. {Anti}-Affinity, or Taint/Toleration
5. Thinking in Kubernetes

# Thanks!

- [github.com/Huang-Wei/shared-loadbalancer](https://github.com/Huang-Wei/shared-loadbalancer)
- Github: @Huang-Wei / @brahmaroutu
- Slack: @Huang-Wei / @srbrahma
- Twitter: @hweicdl / @brahmaroutu

# Backup: CRD Practices

1. Use CRD built-in features
  - a. validation
  - b. shortNames
  - c. additionalPrinterColumns
  - d. controllerRef
  - e. finalizers
2. CRD controller
  - a. kubebuilder
  - b. internal cache
  - c. reconcile upon IndexKey (namespace/name) of a changed object

# Changes in User's View

**Before/Input:** N service yamls (with type LoadBalancer, and src/dst port info)

**Before/Output:** N publicly accessible {ip/hostname, port} pairs

**After/Input:** N custom resource yamls (w/ or w/o src port info)

**After/Output:** 1 publicly accessible {ip/hostname, (random) port} pair

(N is configurable)